

Statistical Inference and Multivariate Analysis (MA324)

LECTURE SLIDES
Topic 03
Multiple Linear Regression

Jan-May 2026

Multiple Linear Regression

- In general, the response (y) may be related to p regressors (input variables/predictors).
- The model

$$y_i = \beta_0 + \beta_1 x_{i1} + \dots + \beta_p x_{ip} + \epsilon_i, i = 1, \dots, n$$

is called **multiple linear regression**. A regression model with p regressors.

- The parameters $\beta_j, j = 0, 1, 2, \dots, p$ are called regression coefficients.
- $\beta_j, j = 0, 1, 2, \dots, p$ represents the change in the average value of the response for a unit change in j^{th} regressor keeping other regressors fixed.

- Data (of sample size n): $(y_i, x_{i1}, x_{i2}, \dots, x_{ip}), i = 1, \dots, n$
- Scalar notation becomes cumbersome – so matrix notation is used
- Let

$$y = \begin{pmatrix} y_1 \\ \vdots \\ y_n \end{pmatrix}, X = \begin{pmatrix} 1 & x_{11} & \dots & x_{1p} \\ \vdots & \vdots & & \vdots \\ 1 & x_{n1} & \dots & x_{np} \end{pmatrix}, \beta = \begin{pmatrix} \beta_0 \\ \vdots \\ \beta_p \end{pmatrix}, \epsilon = \begin{pmatrix} \epsilon_1 \\ \vdots \\ \epsilon_n \end{pmatrix}$$

- Then we can write the model in a more compact form:

$$y_{n \times 1} = X_{n \times (p+1)} \beta_{(p+1) \times 1} + \epsilon_{n \times 1}$$

- X is called the ***design matrix***

Error Assumptions

- As before, ϵ_i 's satisfy
 - $E(\epsilon_i) = 0$ for $i = 1, \dots, n$
 - $Var(\epsilon_i) = \sigma^2$ for $i = 1, \dots, n$
 - $Cov(\epsilon_i, \epsilon_j) = 0$ for all $i \neq j$
- Equivalently, the assumptions can be written as
 - $E(\epsilon) = \mathbf{0}$
 - $Var(\epsilon) = \sigma^2 I_n$

Estimation of Regression Coefficients:

- Consider the objective function

$$\begin{aligned}Q(\beta) &= Q(\beta_0, \beta_1, \dots, \beta_p) \\&= \sum_{i=1}^n \left(y_i - \beta_0 - \beta_1 x_{i1} - \dots - \beta_p x_{ip} \right)^2 \\&= (y - X\beta)^T (y - X\beta) \\&= (y^T y - 2\beta^T X^T y + \beta^T X^T X \beta)\end{aligned}$$

- The LSEs of $\beta_0, \beta_1, \dots, \beta_p$ are $\underset{\beta}{\operatorname{argmin}} Q(\beta)$.

Vector Derivative : Linear Function

- Let x be an $n \times 1$ column vector (the variable).
- Let a be a constant $n \times 1$ column vector.
- The derivative of a scalar-valued linear function with respect to the vector x is

$$\frac{\partial}{\partial x}(a^T x) = a.$$

Vector Derivative: Quadratic Function

- Let A be a constant $n \times n$ matrix.
- The derivative of a quadratic form with respect to the vector x is

$$\frac{\partial}{\partial x} (x^T A x) = (A + A^T) x.$$

- If A is a symmetric matrix, then the above formula simplifies to

$$\frac{\partial}{\partial x} (x^T A x) = 2Ax.$$

Estimation of Model Parameters:

- Differentiating $Q(\beta)$ with respect to β and setting the derivative to zero gives the following *normal equations*:

$$X^T X \beta = X^T y.$$

- The above system of linear equations in β is always consistent.
- If the matrix $X^T X$ is *invertible* (i.e. if X is of *full rank*), then the LSE is unique and given by

$$\hat{\beta} = (X^T X)^{-1} X^T y.$$

Fitted Value:

- The fitted value of response corresponding to regression $x = (1, x_1, \dots, x_p)$ is

$$\hat{y} = \hat{\beta}_0 + \hat{\beta}_1 x_1 + \dots + \hat{\beta}_p x_p$$

- Then $\hat{\mathbf{y}} = (\hat{y}_1, \hat{y}_2, \dots, \hat{y}_n)^T = X\hat{\beta} = X(X^T X)^{-1} X^T \mathbf{y} = H\mathbf{y}$, where, $H = X(X^T X)^{-1} X^T$ is called **hat-matrix**.
- The residuals:

$$\mathbf{e} = \begin{bmatrix} e_1 \\ \vdots \\ e_n \end{bmatrix} = \begin{bmatrix} y_1 - \hat{y}_1 \\ \vdots \\ y_n - \hat{y}_n \end{bmatrix} = \mathbf{y} - \hat{\mathbf{y}} = (I - H)\mathbf{y}$$

Properties of LSE of β

- $\hat{\beta}$ is a linear function of y
- $\hat{\beta}$ is unbiased estimator of β . That is, $E(\hat{\beta}) = \beta$
- $Var(\hat{\beta}) = \sigma^2(X^T X)^{-1}$
- $\hat{\beta}$ is the BLUE of β .

Estimation of Error Variance (σ^2)

- It can be shown that

$$E(SS_{Res}) = (n - p - 1)\sigma^2$$

- Hence, $\hat{\sigma}^2 = \frac{SS_{Res}}{n-p-1} = MS_{Res}$ is an unbiased estimator of σ^2 . Here MS_{Res} is residual mean square.
- Observed value of $\hat{\sigma}^2 = \frac{SS_{Res}}{n-p-1}$ is called **Residual variance**. It's **square root is called residual standard error**.
- $\frac{(n-p-1)MS_{Res}}{\sigma^2} \sim \chi^2_{n-p-1}$.
- Here,

$$SS_{Res} = \sum_{i=1}^n e_i^2 = y^T y - \hat{\beta}_1^T X^T y,$$

Hypothesis Testing: Test for Significance of Regression

- Assumption: $\epsilon \sim N_n(0, \sigma^2 I_n)$.
- Want to test the hypothesis if there is a linear relationship between the response y and any of the regressors $x_1, \dots, x_n$:

$$H_0 : \beta_1 = \beta_2 = \dots = \beta_p = 0 \text{ ag. } H_1 : \beta_j \neq 0 \text{ for atleast one } j = 1, \dots, p$$

- The test statistic is

$$F_0 = \frac{SS_{Reg}/p}{\hat{\sigma}^2} = \frac{SS_{Reg}/p}{SS_{Res}/(n-p-1)} \sim F_{p,n-p-1}, \text{ under } H_0.$$

- Reject H_0 iff $F_0 > F_{p,n-p-1;\alpha}$ (at level α)

Hypothesis Testing: Significance of regression coefficients

- Want to test

$$H_0 : \beta_j = 0 \text{ ag. } H_1 : \beta_j \neq 0$$

- Note that

$$\hat{\beta} \sim N_p \left(\beta, \sigma^2 (X^T X)^{-1} \right) \implies \hat{\beta}_j \sim N \left(\beta_j, \sigma^2 C_{jj} \right),$$

where C_{jj} is the j -th diagonal element of $(X^T X)^{-1}$

- The test statistic is

$$t_0 = \frac{\hat{\beta}_j}{\sqrt{MS_{Res} C_{jj}}} \sim t_{n-p-1} \text{ under } H_0$$

- Reject H_0 iff $|t_0| > t_{n-p-1, \alpha/2}$ (at level α)

Test of contribution of a subset of the regressors



$$y = X\beta + \epsilon \implies y = (X_1, X_2) \begin{pmatrix} \beta_1 \\ \beta_2 \end{pmatrix} + \epsilon$$

- Want to test:

$$H_0 : \beta_2 = 0 \text{ ag. } H_1 : \beta_2 \neq 0$$

- Based on the full model ($y = X\beta + \epsilon$), $\hat{\beta} = (X^T X)^{-1} X^T y$ and $SS_{Reg}(\beta)$ has p degrees of freedom.
- To find the contribution of β_2 , fit the model assuming $\beta_2 = 0$.
- The reduced model is $y = X_1\beta_1 + \epsilon$, $\hat{\beta}_1 = (X_1^T X_1)^{-1} X_1^T y$ and $SS_{Reg}(\beta_1)$ has $p - r$ degrees of freedom. Where r denotes the number of components in β_2

Test of contribution of a subset of the regressors

- $SS_{Reg}(\beta_2 \mid \beta_1) = SS_{Reg}(\beta) - SS_{Reg}(\beta_1)$ can be used as a measure of contribution of β_2
- Note that if β_2 has significant contribution then $SS_{Reg}(\beta_2 \mid \beta_1)$ is large
- Therefore, the test statistic is

$$F_0 = \frac{SS_{Reg}(\beta_2 \mid \beta_1)/r}{\hat{\sigma}^2} = \frac{SS_{Reg}(\beta_2 \mid \beta_1)/r}{SS_{Res}/(n-p-1)} \sim F_{r,n-p-1}, \text{ under } H_0$$

- Reject H_0 iff $F_0 > F_{r,n-p-1;\alpha}$ (at level α)

Confidence intervals for β_j :

- Confidence Interval of individual regression coefficient $\hat{\beta}_j$

- Pivot is

$$\frac{\hat{\beta}_j - \beta_j}{\sqrt{MS_{Res}C_{jj}}} \sim t_{n-p-1},$$

where C_{jj} is the diagonal element of $(X^T X)^{-1}$ matrix.

- A $100(1 - \alpha)\%$ CI for β_j is

$$\left[\hat{\beta}_j \pm t_{n-p-1, \alpha/2} \sqrt{MS_{Res}C_{jj}} \right].$$

Confidence interval for mean response:

- Let $x_0 = \begin{pmatrix} 1 \\ x_{01} \\ \vdots \\ x_{0p} \end{pmatrix}$ be a value of the regressor vector
- The mean response at x_0 is $x_0^T \beta$
- Pivot is $\frac{\hat{y}_0 - x_0^T \beta}{\sqrt{MS_{Res} x_0^T (X^T X)^{-1} x_0}} \sim t_{n-p-1}$
- A $100(1 - \alpha)\%$ CI for $x_0^T \beta$ is

$$\left[\hat{y}_0 \pm t_{n-p-1, \alpha/2} \sqrt{MS_{Res} x_0^T (X^T X)^{-1} x_0} \right]$$

Prediction Interval for New Observation

- Consider a level of regressor x_0 .
- Let y_0 be the corresponding value of the response.
- We want a prediction interval for y_0 .
- Let $\Psi = y_0 - \hat{y}_0 = y_0 - x_0^T \hat{\beta} \sim N(0, \sigma^2(1 + x_0^T (X'X)^{-1} x_0))$.
- A pivot is $\frac{\Psi - 0}{\sqrt{MS_{Res}(1 + x_0^T (X'X)^{-1} x_0)}} \sim t_{n-p-1}$.
- A $100(1 - \alpha)\%$ prediction interval of y_0 is

$$\left[\hat{y}_0 \mp t_{n-p-1; \frac{\alpha}{2}} \sqrt{MS_{Res}(1 + x_0^T (X'X)^{-1} x_0)} \right]$$

Standardized Regression Coefficients

- Difficult to compare regression coefficients. The magnitude of β_j reflects the unit of measurement of regressor x_j .
- For example, $y = 5 + x_1 + 1000x_2$, where y is measured in liters, x_1 in milliliters, and x_2 in liters. Here, $\beta_2 \gg \beta_1$. But the effects of both regressor on y are identical.
- Wayout is to standardized the regressors and response so that they become unit free.

Standardized Regression Coefficients

- A popular approach is as follows:

- Define $W_{ij} = \frac{x_{ij} - \bar{x}_j}{\sqrt{S_{jj}}}$, $i = 1, 2, \dots, n$; $j = 1, 2, \dots, p$.

$$y_i^* = \frac{y_i - \bar{y}}{\sqrt{SS_T}}, i = 1, 2, \dots, n.$$

$$\text{Here } S_{jj} = \sum_{i=1}^n (x_{ij} - \bar{x}_j)^2, j = 1, 2, \dots, p.$$

- $\bar{W}_j = 0$ and $\sqrt{\sum_{i=1}^n (W_{ij} - \bar{W}_j)^2} = \sqrt{\sum_{i=1}^n W_{ij}^2} = 1$
 $\bar{y}^* = 0$ and $\sqrt{\sum_{i=1}^n (y_i^*)^2} = 1$

- In terms of y^* , $W_1, \dots, W_p$, the regression model becomes,

$$y^* = b_1 W_1 + b_2 W_2 + \dots + b_p W_p + \epsilon$$

- In matrix notation: $y^* = Wb + \epsilon$.
 - LSE: $\hat{b} = (W'W)^{-1}W'y^*$.

Standardized Regression Coefficients

- $W'W = \begin{pmatrix} 1 & r_{12} & \cdot & \cdot & \cdot & r_{1p} \\ r_{21} & 1 & \cdot & \cdot & \cdot & r_{2p} \\ \cdot & \cdot & & & & \cdot \\ \cdot & \cdot & & & & \cdot \\ \cdot & \cdot & & & & \cdot \\ r_{p1} & r_{p2} & \cdot & \cdot & \cdot & 1 \end{pmatrix}_{p \times p}$ $W'y^* = \begin{pmatrix} r_{1y} \\ r_{2y} \\ \cdot \\ \cdot \\ \cdot \\ r_{py} \end{pmatrix}_{p \times 1}$

- $r_{ij} = \frac{\sum_{u=1}^n (x_{ui} - \bar{x}_i)(x_{uj} - \bar{x}_j)}{\sqrt{S_{jj}S_{ii}}} = \frac{S_{ij}}{\sqrt{S_{jj}S_{ii}}}$

- $r_{iy} = \frac{\sum_{u=1}^n (x_{ui} - \bar{x}_i)(y_u - \bar{y})}{\sqrt{S_{jj}S_{ii}}} = \frac{S_{jy}}{\sqrt{S_{jj}SS_T}}$

Model Adequacy Checking

- Major assumptions
 - linear relationship
 - Error mean zero
 - Error variance is constant
 - Error are uncorrelated
 - Error are normally distributed
- Gross violation of the assumptions may lead to a totally different model with opposite conclusions.
- We perform the checking using residuals.

- We now define 3 types of residuals.

- Residual: $e_i = y_i - \hat{y}_i$.

- Standardized residual: $d_i = \frac{e_i}{\sqrt{MS_{Res}}}$

- Studentized residual:

$$e = (I - H)\epsilon \Rightarrow \text{Var}(e) = \sigma^2(I - H).$$

$$\text{Var}(e_i) = \sigma^2(1 - h_{ii}) \text{ and } \text{Cov}(e_i, e_j) = -\sigma^2 h_{ij}.$$

It can be shown that $0 \leq h_{ii} \leq 1$.


$$r_i = \frac{e_i}{\sqrt{MS_{Res}(1-h_{ii})}}.$$

Residual Plots: Q-Q Plot (Test for normality)

- The residuals can be assessed for normality using a Q-Q plot. This compares the residuals to “ideal” normal observations.

- We plot the quantiles corresponding to sorted residuals (e_i) against $\Phi^{-1}(\frac{i}{n+1})$ for $i = 1, \dots, n$.


- Four Q-Q plots are:
Normal: ideal;
Lognormal: an example of a skewed distribution;
Cauchy: an example of a long-tailed (platykurtic) distribution;



- **Uniform:** an example of a short-tailed (leptokurtic) distribution
- Source (book): Linear Models with R by Julian J. Faraway

Residual Plots: Plot of residual against Fitted values

- Plot $\hat{y}_i$ vs e_i (or d_i or r_i): (Test of constant variance and non-linear relation)



Residuals vs Fitted plots - the **first** suggests **no change** to the current model while the **second** shows **non-constant variance** and the **third** indicates **some nonlinearity** which should prompt some change in the structural form of the model

Residual Plots

- Plot of residual against regressor
 - Plot x_{ij} vs $e_i \forall j$
 - In previous plot, replace 'fitted' by x_{ij} to find similar interpretation.
 - Repeat it for all the regressors.

● Partial Regression Plot:

- Complete/correct marginal effect
- Let we want to study the marginal effect of x_j on y
- y is regressed on $x_1, \dots, x_k$ except x_j .
- x_j is regressed on $x_1, \dots, x_k$ except x_j .
- Plot y residual, $e_i(y|x_1, \dots, x_{j-1}, x_{j+1}, \dots, x_k)$ against x_j residual, $e_i(x_j|x_1, \dots, x_{j-1}, x_{j+1}, \dots, x_k)$
- For ideal scenario, the partial regression plot should show a linear relationship (straight line with non-zero slope).
- Curvilinear band: indicates higher order terms in x_j or it's transformation.
- Horizontal band: indicates no additional useful information in x_j

The PRESS Statistic

- Delete i^{th} observation. Fit the model on remaining $(n - 1)$ observations. Now, predict y_i . Let the corresponding prediction error be $e_{(i)} = y - \hat{y}_{(i)}$ (PRESS Residual).
- $e_{(i)} = \frac{e_i}{1 - h_{ii}}$.
- Large values of $e_{(i)}$ implies potential influential observations.
- Large difference between e_i and $e_{(i)}$ indicates an observation where model fit is quite well but a model built without that predicts poorly.
- $Var(e_{(i)}) = \frac{\sigma^2}{1 - h_{ii}}$.

- Standardized PRESS residual is $\frac{e_{(i)}}{\sqrt{\text{Var}(e_{(i)})}} = \frac{e_i}{\sqrt{\sigma^2(1-h_{ii})}}$, which is standardized residual if we estimate σ^2 by MS_{Res} .
- $PRESS = \sum_{i=1}^n e_{(i)}^2 = \sum_{i=1}^n \left(\frac{e_i}{1-h_{ii}} \right)^2$.
- PRESS is a measure of how well a regression model perform in predicting new observations.
- R^2 for prediction :
 $R_{prediction}^2 = 1 - \frac{PRESS}{SST}$: gives indication of the prediction capability of the regression model.
- Using PRESS, we may compare model.

Variable Selection

- Criteria for Evaluating Subset Regression Models:

- R^2
- R_{Adj}^2
- Residual Mean Square : $MS_{Res}(p) \equiv R_{Adj}^2$.
- Partial Correlation
- PRESS Statistic

- Techniques:

- All possible Regression.
- Step-wise Type Procedures :
 - Forward Selection : $F = \frac{SS_R(x_2|x_1)}{MS_{Res}(x_1, x_2)}$
 - Backward elimination
 - Step-wise Regression